



Ensuring Kubernetes manifests validity & compliance

A tooling overview

Yann Hamon
CNCF Meetup Berlin
Dec. 18, 2020



About me

- French, German
- Staff Engineer at [@Contentful](#)
- Interested in reliability, security and performance in distributed systems
- [CoreDNS-Nodecache](#), [Kube-secret-syncer](#), [Kubeconform](#)



But first...

Thank you Gareth!

- Author of or contributor to several of the tools presented here (Kubetest, Kubeval, Conftest, Rego rules for Kubesec checks)
- Keeping the resources those tools rely on up-to-date for many years
- Supportive & helpful



Gareth Rushgrove

garethr

Unfollow

...

🔗 800 followers · 11 following · ☆ 257

📁 @snyk

📍 Cambridge

✉ gareth@morethanseven.net

🌐 morethanseven.net



A simple Nginx deployment

Copied from the [Kubernetes documentation](#):

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
```



↓

```
template:
  metadata:
    labels:
      app: nginx
  spec:
    containers:
      - name: nginx
        image: nginx:1.14.2
        ports:
          - containerPort: 80
```



Associated service

I wrote this myself!!

```
apiVersion: v1
kind: service
metadata:
  name: nginx-service
spec:
  ports:
    - port: 80
      protocol: tcp
  selector:
    run: nginx-service
```




Our mission, should you choose to accept it...

*Use available, open-source tooling
to identify all the errors in this manifest*



You will need to...

- Ensure the file is a valid YAML
- ... that is actually a valid Kubernetes manifest
- ... that conforms to your best practices
- ... that would actually apply without errors





YAML conformity

YAML validators / linters - [yq](#), [yamllint](#)



Validating YAML - yq

- Multiple ways of doing so: yq, rq, python one-liners, ...
- yq is a simple YAML validator

```
$ yq validate fixtures/nginx-service.yml
```

```
Error: yaml: line 7: mapping values are not allowed in this context
```



Validating / linting YAML: Yamllint

- Will also detect trailing spaces, inconsistent alignment, duplicate keys...
- Can recursively test folders

```
$ yamllint -d relaxed fixtures/  
fixtures/nginx-service.yml  
5:6      error      trailing spaces (trailing-spaces)  
7:9      error      syntax error: mapping values are not allowed here (syntax)
```



Yamllint

Pros:

- Very powerful and flexible
- Integrate with vim, emacs, more

Cons:

- YAML errors also covered by higher-level tools
- Not very fast
- Errors are mostly cosmetic



Fixed manifest

One validation error, one trailing whitespace

```
apiVersion: v1
kind: service
metadata:
  name: nginx-service
spec:
  ports:
    - port: 80
      protocol: tcp
  selector:
    run: nginx-service
```



Learnings

- Most higher-level tools will also validate YAML
- While consistent formatting is nice...
- Failing tests for whitespaces can be frustrating
- Use YAML linters in your editor

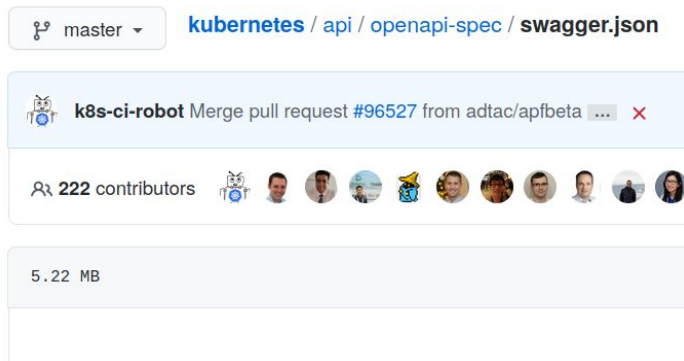


Kubernetes manifest conformity

Kubernetes validators - [Kubeval](#), [Kubeconform](#), [Kubectrl](#)

Manifests should conform to K8S schemas

- Kubernetes publishes a Swagger file to describe its API
- Describes the Kinds of manifests, required/optional properties, types...
- Conforming to the schema is necessary but not sufficient for a manifest to be accepted by the API Server
- Especially important when upgrading to a new version of Kubernetes!





Validating a K8S manifest - Kubeval

```
$ kubeval --strict fixtures/nginx-deployment.yml fixtures/nginx-service.yml
PASS - fixtures/nginx-deployment.yml contains a valid Deployment
WARN - fixtures/nginx-service.yml contains an invalid service - kind: kind must be
one of the following: "Service"
```

```
$ echo $?
1
```

Simple enough. How does it work?

Kubeval's workflow

master [kubernetes / api / openapi-spec / swagger.json](#)

k8s-ci-robot Merge pull request #96527 from adtac/apfbeta ... x

222 contributors

5.22 MB

instrumenta / [kubernetes-json-schema](#)

Schemas for every version of every object in every version of Kubernetes

[kubernetesjsonschema.dev](#)

View license

☆ 116 stars 🍴 31 forks

instrumenta / [openapi2jsonschema](#)

Convert OpenAPI definitions into JSON schemas for all types in the API

View license

☆ 122 stars 🍴 40 forks

Kubeval





Kubeval

Pros:

- Widely used
- Can iterate recursively over folders
- Can validate against different versions of Kubernetes

Cons:

- No support for Custom Resources*
- Some not-so-easy-to-fix bugs

validation false positives if input is stdin #133



jkroepke opened this issue on Jun 9, 2019 · 1 comment

-o json contains non-json formatted output #238



timgm opened this issue on May 27 · 0 comments

Invalid deployment gives warning --> returns code 1 #233



audunsolemdal opened this issue on May 4 · 4 comments



Custom Resources support

- Welcome to the world of Kubernetes Operators
- Contentful uses the Prometheus, Jaeger, Kube-secret-syncer, and other operators... Custom Resources are everywhere
- One near-incident was due to missing validation of those files



Introducing Kubeconform

```
$ kubeconform -summary nginx-service.yml
```

```
Summary: 1 resource found in 1 file - Valid: 1, Invalid: 0, Errors: 0, Skipped: 0
```

```
$ echo $?
```

```
0
```



Kubeconform's CRD support

```
$ ./scripts/openapi2jsonschema.py \  
https://raw.githubusercontent.com/aws/amazon-sagemaker-operator-for-k8s/master/config/crd/bases/sagemaker.aws.amazon.com\_trainingjobs.yaml
```

```
JSON schema written to trainingjob_v1.json
```

```
$ ./bin/kubeconform -schema-location './{{ .ResourceKind }}{{ .ResourceAPIVersion }}.json' fixtures/trainingjob-customresource.yaml
```

```
$ echo $?
```

```
0
```



Kubeconform

Pros:

- Feature-parity with Kubeval, same test-suite
- Fixes some outstanding Kubeval bugs
- Some performance improvements
- Flexible support for Custom Resources Schemas

Cons:

- Smaller community, not as battle-tested



Kubectl

```
$ kubectl apply --validate=true --dry-run=client -f fixtures/nginx-service.yml  
service/nginx-service created (dry run)
```

```
$ echo $?  
0
```

... the capitalisation error is not detected?



Kubecttl #2

This time with a broken deployment...

```
$ kubecttl apply --validate=true --dry-run=client -f fixtures/nginx-deployment-invalid.yml
error: error validating "fixtures/nginx-deployment-invalid.yml": error validating data:
ValidationError(Deployment.spec.template.spec.containers[0].ports[0].containerPort):
invalid type for io.k8s.api.core.v1.ContainerPort.containerPort: got "string", expected
"integer"; if you choose to ignore these errors, turn validation off with
--validate=false
```

```
$ echo $?
```

```
1
```




Kubectl

Pros:

- “Kubernetes upstream” project

Cons:

- Requires a connection to the Kubernetes cluster (even in client-side validation mode? [#991](#))
- Unclear what is validated / poorly documented feature



Fixed manifest

“Service” needs to be capitalized

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  ports:
    - port: 80
      protocol: tcp
  selector:
    run: nginx-service
```



Learnings

- Use ~~Kubeval~~ Kubeconform to validate your manifests
- Similar functionality seems to be arriving in Kubectl
 - but poorly documented as of 2020
 - some validation errors not caught



Enforcing best practices / compliance

Testing Frameworks - [Conftest](#), [Kubesecc](#)



Best-practices, compliance

- “All deployments should have resource requests set”
- “All manifests should have a namespace explicitly defined”
- “No containers should run as privileged”
- “Resource names shall be lower-case”
- “Processes in containers should not run as root”
- [...]



Conftest, your manifests' policy framework

- Policies written in Rego - a language inspired by Datalog, itself derived from.. Prolog
- A declarative language - “Simpler & more concise” (.. supposedly)



“All manifests should have a namespace explicitly defined”

```
$ conftest test -p fixtures/policies/namespace.rego fixtures/nginx-deployment.yml  
fixtures/nginx-service.yml  
FAIL - fixtures/nginx-service.yml - no namespace set  
FAIL - fixtures/nginx-deployment.yml - no namespace set
```

```
2 tests, 0 passed, 0 warnings, 2 failures
```

```
$ echo $?
```

```
1
```



"All manifests should have a namespace explicitly defined"

```
package main
```

```
kinds_to_skip = {  
    "Apiservice",  
    "Clusterrole",  
    [...]  
}
```

```
deny[msg] {  
    not kinds_to_skip[lower(input.kind)]  
    not input.metadata.namespace  
    msg = "no namespace set"  
}
```




Confest

Pros:

- Powerful language
- Same language used for Admission controllers
- Quality of [documentation](#)
- A lot of [existing rules](#)
- Large community

Cons:

- Rego can be unfamiliar and “one more thing” to pick up



Security testing - Kubesec

Scores your manifests using a list of built-in rules

```
$ kubesec scan fixtures/nginx-*.yaml
[...]
```

```
{
  "selector": "containers[] .securityContext .readOnlyRootFilesystem == true",
  "reason": "An immutable root filesystem can prevent malicious binaries being added to PATH and
increase attack cost",
  "points": 1
}, {
  "selector": "containers[] .securityContext .runAsNonRoot == true",
  "reason": "Force the running image to run as a non-root user to ensure least privilege",
  "points": 1
},
[...]
```



Kubesecc

Pros:

- Insightful rules
- Each rule is well-documented on their website with a description of potential impact ([example](#))

Cons:

- Limited scope
- Check-out [kubesecc policies in Rego](#)



Honourable mention: Kube-score

- Similar approach as Kubesec - scores manifests according to built-in rules
- Not extensible
- Good, documented [list of checks](#)

Honourable mention: Kube-linter

- Similar to kube-score, but will report errors instead of scores
- Not extensible
- In early stage of development



Honourable mention: Kubetest

- Testing Framework with tests written in Skylark, a subset of Python
- Archived in favour of Conftest



Fixed manifest

Set the namespace explicitly

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
  namespace: default
spec:
  ports:
    - port: 80
      protocol: tcp
  selector:
    run: nginx-service
```



Learnings

- Conftest & OPA is the de-facto standard for policy-based testing
- Most people are not familiar with Rego, but it is not too hard to pick up
- It should be part of your CI!



Testing groups of resources

Featuring Conftest



Testing groups of resources

- “A deployment should specify a number of replicas, unless it has a matching HorizontalPodAutoscaling rule”
- “A service's 'selector' should have a matching deployment”
- “The configmap referenced by a deployment should exist”

All this is made possible by “conftest test --combine”

Warning: runtime complexity increases with the number of files tested



"A service's 'selector' should have a matching deployment"

```
package main

deployment[[input, file]] {
    input[deploymentfile].kind == "Deployment"
    input[deploymentfile].spec.selector.matchLabels.app == input[file].spec.selector.run
}

deny[msg] {
    input[file].kind == "Service"
    not deployment[[input, file]]

    msg := sprintf("Service %v points to non-existing deployment %v" ,
[input[file].metadata.name, input[file].spec.selector.run])
}
```



Now...

```
apiVersion: v1
kind: Service
metadata:
  name: nginx
  namespace: default
spec:
  ports:
    - port: 80
      protocol: TCP
  selector:
    run: nginx-service
```



Does the manifest apply without errors?

Featuring [Kubectrl](#)



Kubernetes & kubectl validate differently

Kubernetes' API, Controllers and Validation Webhooks add an extra layer of validation

```
$ kubectl apply --validate=true --dry-run=client -f fixtures/nginx-service.yml  
service/nginx configured (dry run)
```

```
$ kubectl apply --validate=true --dry-run=server -f fixtures/nginx-service.yml  
The Service "nginx" is invalid: spec.ports[0].protocol: Unsupported value: "tcp":  
supported values: "SCTP", "TCP", "UDP"
```

```
$ echo $?  
1
```



Fixed manifest

Protocol should be upper-case

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
  namespace: default
spec:
  ports:
    - port: 80
      protocol: TCP
  selector:
    run: nginx-service
```



Conclusion



Final manifest

```
apiVersion: v1
kind: Service # Kubeconform, yamllint
metadata:
  name: nginx
  namespace: default # Conftest
spec:
  Ports: # yq, yamllint
  - port: 80
    protocol: TCP # Kubectl (dry-run=server)
  selector:
    run: nginx-service # Conftest (across multiple files)
```




You should probably use...

- A YAML linter integrated with your editor
- Kubeconform to ensure your files are valid Kubernetes manifests
- Conftest with a library of policies for additional validation, best-practices and security compliance



Errors in manifests are easy to overlook

.. you should test your manifests!

- It is easy
- Save time by detecting errors earlier
- Kubernetes will not detect all mistakes - some might lead to operational problems
- Enforce best practices
- Ensure security compliance



Questions?

